

 Dahir-ke / dsc_phase3_project  

<> Code ⚙ Issues ⚡ Pull requests ⚡ Actions Projects 📄 Wiki 🛡 Security 🔍 Insights ⚙ Settings

 Watch 0    

star 0 stars fork 0 forks watching 0 Branches Activity Tags

Public repository

 main  1 Branch  0 Tags   Go to file  Go to file  Add file  

 Dahir-ke updated projected 2fb6ab8 · 6 minutes ago 

 Images	Second commit	2 days ago
 Presentation.pdf	Project updated	17 minutes ago
 README.md	updated README.md	14 minutes ago
 Test_set_values.csv	first commit	3 days ago
 Training_set_labels.csv	combined data sets	2 days ago
 phase3_project.ipynb	updated projected	6 minutes ago
 phase3_project.pdf	updated projected	6 minutes ago

 README  

Water Well Status Prediction in Tanzania

Project Overview

Access to clean and reliable water is a cornerstone of public health and economic development in Tanzania. However, thousands of water wells across the country suffer from breakdowns, poor maintenance, or complete failure. This project uses data from the **Tanzanian Water Point Mapping** initiative to build a predictive model that determines the operational status of a water well.

Primary Stakeholders:

- Ministry of Water (Tanzania)
- Water-focused NGOs (e.g., WaterAid, World Vision)

Goal:

To maximise the number of functional water wells by identifying which water points are most likely to fail, enabling proactive maintenance and smarter investment decisions.

Dataset

- Source: [Tanzanian Water Point Mapping](#)
- Size: 59,400 rows, 41 columns
- Target variable: status_group – three classes:
 - functional (encoded as 2)
 - functional needs repair (encoded as 1)

- non functional (encoded as 0)
- Features include:
 - Technical: pump type, extraction type, construction year, amount_tsh (water yield)
 - Environmental: GPS coordinates, elevation, basin, region
 - Management: funder, installer, payment type, permit, public meeting, scheme management
- Missing data & invalid values:
 - Categorical columns with low missing percentages imputed with mode or "unknown".
 - Numeric columns contained invalid zeros (e.g., gps_height = 0, longitude = 0, construction_year = 0). These were replaced using domain-aware imputation (non-zero medians, regional means for coordinates).
 - scheme_name (~49% missing) was dropped.

Methodology

1. Data Preprocessing & Feature Engineering

Step	Description
Imputation	Zero values in amount_tsh, gps_height, population, construction_year replaced with the median of non-zero entries. Geographic coordinates imputed by regional means.
Noise addition	Added random noise (± 5 years) to construction_year to break an artificial spike at 2000 (caused by mass imputation of zeros).
Feature engineering	<pre>well_age = 2013 - construction_year amount_tsh_log = log1p(amount_tsh) population_log = log1p(population) gps_height_bin - binned elevation into low, medium, high.</pre>
Encoding	High-cardinality columns (funder, installer, subvillage, lga, ward) → Label Encoding. Low/medium-cardinality columns → One-Hot Encoding with drop_first=True. Target encoded ordinally: non functional = 0, needs repair = 1, functional = 2.
Train/Test split	75% training, 25% testing, stratified by target.

2. Modeling Approach

Model	Rationale	Key Outcome
Logistic Regression	Interpretable baseline.	ROC-AUC = 0.6846; recall for class 1 = 0.00 → unusable for business goal.
Decision Tree	Captures non-linear patterns; human-readable rules.	ROC-AUC = 0.7771; recall for class 1 = 0.07 → still misses most at-risk wells.
Random Forest	Ensemble reduces variance; handles complex interactions.	ROC-AUC = 0.8644; recall for class 1 = 0.05 → excellent overall, but fails on minority class.

3. Addressing Class Imbalance

Three Random Forest variants were compared:

Model	ROC-AUC (OvR)	Accuracy	Weighted F1	Recall (class 1)
Base RF (tuned)	0.8654	0.75	0.72	0.05
RF + class_weight='balanced'	0.8733	0.73	0.75	0.66
RF + SMOTE	0.8651	0.76	0.76	0.40

Hyperparameter tuning: RandomizedSearchCV on a 30% subset, then final model trained on full dataset.

Final model recommendation: Random Forest with `class_weight='balanced'` – it balances strong overall performance with a 66% recall of wells that need repair (compared to 5% in the base model).

4. Feature Importance

Top predictors of well failure (from the final Random Forest):

1. Water quantity (`quantity`) – dry or insufficient wells are 79%/52% likely to be non-functional.
2. Well age (`well_age`) – failure risk doubles after 20 years.
3. Payment type (`payment_type`) – “never pay” wells fail twice as often.
4. Elevation (`gps_height`) – higher elevation correlates with better functionality.
5. Pump mechanics (`extraction_type`) – some designs (gravity, Afridev) are significantly more reliable.

Results & Business Impact

- Final ROC-AUC: 0.8733 – surpasses the project target of 0.85.
- Minority class recall (Needs Repair): 0.66 – the model now identifies **two out of every three** at-risk wells, enabling **proactive maintenance** before complete failure.
- Functional wells are predicted with **94% recall** – authorities can trust the model when it declares a well operational.

Before this model: maintenance was purely reactive – wells were fixed only after communities called with a crisis.

Now: the Ministry receives a monthly risk-score list and can schedule inspections for the highest-risk wells during routine trips – **no extra travel costs, just smarter routing.**

368 wells correctly identified as “needs repair” – 368 communities that will **not** lose water access. That is the real impact.

Actionable Recommendations (Evidence-Based)

Based on the model’s insights, we propose three immediate, no-regret moves:

1. Target “Dry” & “Insufficient” Wells First
 - Wells flagged as `dry` are 79% non-functional; `insufficient` are 52% non-functional.
 - Action: Deploy rapid response teams to every such well within 7 days.
2. Phase Out Wells Older Than 25 Years
 - Wells built before 1995 are 2.3× more likely to be non-functional.
 - Action: Create a 5-year phased replacement plan for all pre-1995 wells (~12,000 wells).
3. Replace “Never Pay” with Community Contributions
 - Wells with **no payment system** fail 2× more often than those with any fee.
 - Action: Convert the 25,000+ wells currently marked “never pay” to a low-cost community maintenance model (e.g. pay-per-bucket or monthly fee).

How to Run the Notebook

1. Clone the repository

```
git clone https://github.com/Dahir-ke/dsc_phase3_project.git
cd <repository-files.csv>
```





Releases

No releases published

[Create a new release](#)

Packages

No packages published

[Publish your first package](#)

Languages

- Jupyter Notebook 100.0%